
BudgetFlow: Value-Aware Budget Governance for Agent Tasks

Fengde Ding¹

fengde_ding@uic.edu.cn

¹University of International Relations

Abstract

Multi-step AI agent tasks consume variable amounts of model budget, and the tasks in a batch often have unequal value. Existing model-routing work primarily decides which model should answer one request or how one task should spend its own budget. We study *Batch-Level Task Budget Allocation*: a task batch shares one hard budget, and the system must allocate model opportunities across tasks before the budget is exhausted.

We introduce BudgetFlow, a value-aware budget governance framework for task batches. Each task exposes four interfaces: Task Value, Estimated Token Demand, Model Fit, and a Verifier. BudgetFlow uses the first three signals to allocate budget and model tiers, while the Verifier settles outcomes. We evaluate BudgetFlow on SWE-bench-style verified bug-fixing tasks under fixed task sets and shared hard budgets. The paper reports three primary metrics: Resolved Rate, Total Resolved Value (TRV), and TRV per Dollar.

Across one 4-policy run and three 5-policy runs over 30-task batches, BudgetFlow wins the latest audited run on all three primary metrics, wins a 4-policy run on TRV and TRV per Dollar, and closely approaches the strong-model-only boundary in two runs where the strong model is unusually cost-effective. Sensitivity analyses over Task Value, budget cap, and KV-cache cost discount show how the cost-value frontier changes under different operating conditions. These results support BudgetFlow as a concrete formulation of value-aware task-batch budgeting, with SWE-bench as a verified coding instantiation.

1. Introduction

AI agents increasingly execute multi-step tasks: they inspect context, call tools, edit artifacts, run checks, and decide whether another model call is worth the cost. These tasks consume budget unevenly. One task may finish after a short reasoning loop; another may require repeated search, repair, and verification. In a batch, this variation means that a fixed model budget can be exhausted before every task receives the same opportunity to succeed.

A central issue is that tasks rarely have equal value. A support ticket with an imminent service-level agreement, a security patch, a revenue-critical spreadsheet fix, and a routine rewrite should receive different budget treatment when they are submitted in the same batch. Once model calls become a scarce shared resource, the system needs an allocation policy in addition to a model router.

This paper studies *Batch-Level Task Budget Allocation*. The general setting has many tasks, one shared hard budget, model tiers with different costs and capabilities, and an outcome rule that decides whether a task is resolved. The objective is to allocate scarce model opportunities so the batch produces more verified task value under the same resource constraint.

The operational decision is simple: which task should receive the next scarce model opportunity, under the same shared budget?

BudgetFlow is a budget governance framework for this setting. It treats each task as exposing four interfaces:

- **Task Value:** the pre-registered value gained if the task is resolved;
- **Estimated Token Demand:** the pre-run estimate of token/runway demand;
- **Model Fit:** evidence about how well each model tier fits the task;
- **Verifier:** the outcome rule that maps a completed attempt to resolved or unresolved.

The first three interfaces drive allocation decisions. The Verifier settles outcomes and defines the objective. This separation is the core interface. The policy can change how it routes, caps, assigns model opportunities, or stops tasks, while the outcome contract remains explicit.

We instantiate this formulation on SWE-bench-style verified coding tasks. SWE-bench is useful here because it gives

a strict task-batch testbed: each task is a real bug-fixing issue, each attempt has measurable model spend, and each outcome is settled by a test-based harness. The benchmark is the evidence domain, while the allocation interface is the research object.

The paper uses three primary metrics. **Resolved Rate** anchors the evaluation to standard verified success. **Total Resolved Value** (TRV) sums pre-registered Task Value over resolved tasks. **TRV per Dollar** measures value efficiency under the shared hard budget.

Our results are boundary-aware. BudgetFlow wins the latest audited 5-policy, 30-task run on Resolved Rate, TRV, and TRV per Dollar. It wins a 4-policy run on value and value efficiency while resolving one fewer task than the strong-model-only baseline. The strong-model-only policy leads in conditions where the strong model itself forms a strong cost-value boundary. These outcomes are reported together because the scientific object is the cost-value frontier across operating conditions.

This paper makes three contributions:

1. We formulate Batch-Level Task Budget Allocation with pre-registered Task Value, Estimated Token Demand, Model Fit, a shared hard budget, and a Verifier-settled objective.
2. We present BudgetFlow, a task-batch budget governance framework that allocates model opportunities using value, estimated demand, model fit, and remaining budget.
3. We evaluate BudgetFlow against cheap-model-only, strong-model-only, learned task-router, and budget-only baselines, reporting Resolved Rate, TRV, TRV per Dollar, and sensitivity curves that expose operating conditions.

2. Related Work

Per-query routing and cascading. RouteLLM learns per-query routers from preference data (Ong et al., 2025). Cascade Routing studies model routing and cascading for individual queries (Dekoninck et al., 2025). RouteNLP targets enterprise query routing under quality constraints (Guo et al., 2026). UCCI calibrates uncertainty for cost-optimal cascade routing (Kotte, 2026). These works make model selection a first-class cost-quality problem. The unit is typically one prompt or one request. BudgetFlow operates at the task-batch layer: many multi-step tasks share one hard budget, and the policy allocates scarce model opportunities across tasks with different values.

Within-task budgeted agents. INTENT studies budget-constrained tool use inside a single agent task (Liu et al.,

2026). This is a complementary layer to Batch-Level Task Budget Allocation. A within-task policy decides whether one task should spend another tool or model call. BudgetFlow decides how one budget should be shared across many tasks before the batch budget is exhausted.

Coding-agent evaluation and cost-aware runtimes. SWE-bench provides the verified coding-task substrate used by many agent evaluations (Jimenez et al., 2024). OpenSquilla is a token-efficient agent runtime with model routing, diagnostics, replay, and cost-per-success framing (OpenSquilla contributors, 2026). Claw-SWE-Bench studies harnesses, adapters, prompts, patch extraction, leakage cleanup, and cost reporting for SWE-bench-style coding-agent evaluation (Zheng et al., 2026). These systems sharpen the evaluation requirements for agent tasks: cost, harness behavior, and outcome verification must be explicit.

Serving substrates. vLLM improves LLM serving through PagedAttention and KV-cache management (Kwon et al., 2023); SGLang accelerates structured language-model programs with runtime optimizations such as KV-cache reuse (Zheng et al., 2024); Dynamo exposes agentic inference concerns including cache locality and scheduling (NVIDIA, 2026). These systems provide execution substrates. BudgetFlow provides task-level value and budget signals that can inform priority, continuation, model opportunity, and cache-sticky scheduling in future serving-aware systems.

3. Problem Formulation

Let a task batch contain tasks $\mathcal{T} = \{1, \dots, n\}$ and let B denote the shared hard budget. Each task i has a pre-registered value $v_i > 0$, an estimated token demand $d_i > 0$, and model-fit evidence $m_{i,k}$ for each model tier k . A policy π chooses a sequence of actions that spend budget on model calls and task attempts. Let $c_i(\pi)$ be the spend on task i , and let

$$\sum_{i=1}^n c_i(\pi) \leq B$$

be the shared hard-budget constraint.

The budget is shared at the batch level. A policy can spend more on one task only by leaving less budget for later tasks in the same batch. This coupling is what separates the problem from independent per-task routing: a locally reasonable strong-model call can be globally poor if it consumes budget needed by a higher-value task.

Each task has a Verifier V_i that maps the final task evidence to a binary outcome:

$$r_i(\pi) = V_i(\pi, i) \in \{0, 1\}.$$

The primary value objective is Total Resolved Value:

$$\text{TRV}(\pi) = \sum_{i=1}^n v_i r_i(\pi).$$

We report three headline metrics:

$$\text{ResolvedRate}(\pi) = \frac{1}{n} \sum_{i=1}^n r_i(\pi),$$

$$\text{TRVperDollar}(\pi) = \frac{\text{TRV}(\pi)}{\sum_i c_i(\pi)}.$$

3.1. Core Terms and Metric Reading

Table 1 gives the compact reading guide used throughout the paper. The metrics are intentionally simple: success is verifier-settled, value is pre-registered, and efficiency is value divided by spend. This makes the frontier interpretable through standard verified outcomes and a pre-registered value ledger.

Term	Reading
Batch-Level Task Budget Allocation	A task batch shares one hard budget; a 30-task run spends at most \$9.95 total.
Task Value	Pre-registered value gained when the task passes the Verifier.
Estimated Token Demand	Pre-run token/runway estimate used for cap planning.
Model Fit	Evidence about which model tier is likely to help this task.
Verifier	Domain outcome rule; SWE-bench uses a test-based harness.
Resolved Rate	Verifier-passed tasks divided by all tasks; 16/30 means 53.3%.
Total Resolved Value (TRV)	Sum of Task Value over verifier-passed tasks.
TRV per Dollar	TRV divided by total spend; TRV 18.0 at \$9.95 gives 1.81 TRV/\$.

Table 1. Core terms and metrics. The examples show how to read the reported numbers.

The Verifier is an outcome settlement rule. The policy uses Task Value, Estimated Token Demand, and Model Fit for allocation; the Verifier settles the final outcome after execution. In our SWE-bench instantiation, V_i is the test-based harness. This gives a strict external success signal while preserving the task-batch allocation structure.

The four interfaces have different portability requirements. Task Value must be fixed before execution. Estimated Token Demand is a pre-run planning estimate; final spend is measured after execution. Model Fit can come from catalog priors, historical outcomes, or calibration runs. The Verifier must be a trusted outcome rule for the domain. Their concrete calibration changes by domain, but the allocation problem consumes the same interface.

4. BudgetFlow

BudgetFlow maintains a shared budget ledger across the task batch. For each task, it estimates whether the next model opportunity should use the cheap model, the strong model, or defer additional spend. The ledger is hard: the remaining batch budget gates each planned action.

The policy uses three allocation inputs:

- v_i , the Task Value;
- d_i , the Estimated Token Demand;
- $m_{i,k}$, the Model Fit for tier k .

These inputs enter distinct decisions. Task Value prioritizes tasks that create more value if verified. Estimated Token Demand becomes runway and cap planning: the policy estimates how much budget a task may need before it can produce a useful artifact. Model Fit estimates whether the strong model is likely to improve the outcome relative to the cheap model. Remaining budget controls affordability and stop/defer decisions.

Let h denote budget pressure, with higher h indicating less remaining headroom. Let k_c be the cheap model and k_s be the strong model. A simplified task-start score for giving task i a strong-model opportunity is

$$S_i = \frac{v_i \cdot \max(m_{i,k_s} - m_{i,k_c}, 0) \cdot g(d_i)}{\max(\hat{c}_{i,k_s} - \hat{c}_{i,k_c}, \epsilon)} \cdot q(h),$$

where $\hat{c}_{i,k}$ is the estimated cost for tier k , $g(d_i)$ scales the opportunity by estimated runway, and $q(h)$ reduces strong-model access as the shared budget tightens. The implementation uses this score with affordability checks, learned task-router priors, and hard-budget gates. The formula captures the core allocation principle: strong-model budget is most valuable when a task has high value, sufficient expected model-fit gain, enough estimated runway, and enough remaining budget.

BudgetFlow separates allocation from outcome settlement. It uses Task Value, Estimated Token Demand, Model Fit, and the budget ledger to choose model opportunities. The Verifier evaluates the resulting patch or artifact after the task attempt. This design keeps the value objective auditable: the policy can explain why a task received budget, and the verifier independently determines whether that spend produced resolved value.

Allocation decisions. At task start, BudgetFlow assigns a route and an initial cap from value, estimated runway, model fit, and remaining budget. During execution, it can assign a strong-model opportunity when the strong model has expected upside and sufficient budget remains. It can

Interface	SWE-bench instantiation	Other task-batch instantiations
Task Value	pre-registered criticality value	business priority, user tier, task owner value
Estimated Token Demand	pre-run token/runway estimate from task metadata	expected context size, tool steps, artifact size
Model Fit Verifier	tier evidence from prior runs and catalog priors test-based harness	domain-specific success prior by model tier human review, rule check, delayed business outcome

Table 2. Domain interfaces for Batch-Level Task Budget Allocation. The paper evaluates the SWE-bench instantiation.

Algorithm 1 BudgetFlow batch policy sketch

- 1: **Initialize** remaining budget $R \leftarrow B$.
- 2: **for** task i in the batch **do**
- 3: **Read** $(v_i, d_i, m_{i,k_c}, m_{i,k_s})$ and current R .
- 4: **Set** an initial task cap from Task Value, Estimated Token Demand, and R .
- 5: **Choose** cheap or strong model from Model Fit, value upside, and affordability.
- 6: **Execute** the task while updating spend and remaining budget.
- 7: **Continue, stop, or defer** using value upside and budget pressure.
- 8: **Settle** the final artifact with V_i and record r_i .
- 9: **end for**
- 10: **Report** Resolved Rate, TRV, TRV per Dollar, and total spend.

stop or defer when additional spend has low expected value or would threaten the shared budget. These decisions are task-level in the current paper; finer stage or segment routing is left for future work.

Algorithm 1 is the implementation-facing view of the framework. The policy reads the three allocation interfaces before and during execution, while the Verifier is applied after the task attempt. This keeps routing, cap planning, and strong-model opportunity decisions auditable under the same value ledger used for evaluation.

5. Experimental Setup

Task set. We evaluate on fixed 30-task SWE-bench-style bug-fixing batches. Each task requires an agent to inspect a repository, edit code, and produce a patch. Outcomes are scored by the local harness and cross-checked with exported official predictions where available. Task values are pre-registered before execution.

Policies. Table 3 compares the five policies by the inputs they use. All policies run under the same shared hard budget and the same verifier. The key contrast is whether Task Value enters allocation of scarce model opportunities.

The 4-policy run uses the four policies available in that experimental line. All mainline policies use the same task order, budget protocol, model-tier catalog, and scoring path inside a run.

Metrics. The main tables report Resolved Rate, TRV, and TRV per Dollar, as defined in Table 1. Resolved Rate is written as both count and percentage. Spend is included as budget context because TRV per Dollar depends on it.

Sensitivity analyses. The main paid runs show observed frontier points. The sensitivity analyses ask how that frontier changes when the value ledger, the budget cap, or the serving cost model changes while the completed outcomes stay fixed. Task Value sensitivity re-scores fixed outcomes under alternate value profiles. Budget sensitivity replays the observed task order under different shared budget caps. KV Cache Cost-Discount sensitivity re-costs repeated input tokens under explicit cache-discount assumptions. Together, these diagnostics separate outcome quality from the accounting assumptions that determine cost-value efficiency.

6. Results

6.1. Main Paid Evidence

Table 4 summarizes four paid task-batch runs. Red entries mark BudgetFlow frontier signals; blue entries mark the strongest control when a control defines the boundary. The **latest audited 5-policy run is the strongest positive signal**: BudgetFlow resolves 16/30 tasks, reaches TRV 18.00, and achieves 1.81 TRV per Dollar. The next-best learned task-router resolves 15/30 tasks, reaches TRV 17.00, and achieves 1.71 TRV per Dollar.

The runs show three operating conditions. First, value-aware allocation can protect high-value tasks and **improve TRV and TRV per Dollar** even when it resolves one fewer task than strong-model-only. Second, **a strong model can be the best boundary** when it is both capable and cost-effective under the cap. Third, with the latest audited policy, BudgetFlow **improves all three primary metrics** against the five-policy table.

These outcomes should be read together. The 4-policy run

BudgetFlow: Value-Aware Budget Governance for Agent Tasks

Policy	Same budget	Task Value	Runway/cap planning	Model choice	Same verifier
cheap-model-only	yes			fixed cheap model	yes
strong-model-only	yes			fixed strong model	yes
learned task-router	yes			value-blind learned choice	yes
budget-only	yes		yes	budget-pressure choice	yes
BudgetFlow	yes	yes	yes	value-aware choice	yes

Table 3. Policy inputs and shared outcome contract. Runway/cap planning is the decision-facing use of Estimated Token Demand.

Run	Leading policy	Resolved Rate	Spend	TRV	TRV per Dollar
4-policy value-aware allocation case	BudgetFlow	14/30 (46.7%)	\$7.52	18.50	2.46
	strong-model-only	15/30 (50.0%)	\$9.46	18.00	1.90
5-policy strong-model boundary case	strong-model-only	17/30 (56.7%)	\$9.95	20.00	2.01
	BudgetFlow	15/30 (50.0%)	\$9.95	17.50	1.76
5-policy learned-router stress case	strong-model-only	18/30 (60.0%)	\$9.88	19.00	1.92
	BudgetFlow	17/30 (56.7%)	\$9.95	18.50	1.86
Latest audited 5-policy positive case	BudgetFlow	16/30 (53.3%)	\$9.95	18.00	1.81
	learned task-router	15/30 (50.0%)	\$9.95	17.00	1.71
	strong-model-only	15/30 (50.0%)	\$9.95	16.50	1.66

Table 4. Main paid evidence. Each row reports the leader for one 30-task batch. Resolved Rate, TRV, and TRV per Dollar are the three primary metrics. Red bold entries mark BudgetFlow frontier signals; blue bold entries mark the strongest control in boundary runs.

Policy	Resolved Rate	TRV	TRV/\$
BudgetFlow	16/30 (53.3%)	18.00	1.81
learned task-router	15/30 (50.0%)	17.00	1.71
strong-model-only	15/30 (50.0%)	16.50	1.66
budget-only	12/30 (40.0%)	13.00	1.31
cheap-model-only	11/30 (36.7%)	12.00	1.21

Table 5. Latest audited 5-policy run over 30 tasks.

is the clearest value-aware allocation case: raw resolved count alone favors the strong-model-only baseline, but the resolved tasks under BudgetFlow **carry more value and cost less**. The two 5-policy boundary runs show **pressure from strong controls**. The latest audited 5-policy run shows the desired frontier movement: BudgetFlow **leads on standard resolution, value, and value per dollar** under the same cap.

6.2. Full Latest 5-Policy Table

Table 5 gives the full latest audited run. **BudgetFlow leads all three primary metrics**. This table is the main point comparison; the next section adds sensitivity curves to show how the ordering changes under alternate operating conditions.

6.3. Task-Level Model Advantage

The latest run contains real model-tier heterogeneity. Among tasks where both cheap and strong model rows consumed model budget, there are **six cheap-model-favorable passes, five strong-model-favorable passes, and seven comparable failures under both**. This explains why all-

simple policies are useful controls: the strong model often helps, but the task batch also contains tasks where cheap-model execution is the better use of budget.

This task-level heterogeneity motivates selective strong-model use. When strong-model upside is concentrated on a subset of tasks, the allocation policy should preserve strong-model opportunities for those tasks while spending less on tasks where the cheap model is sufficient or both tiers are unlikely to pass.

7. Sensitivity Analysis

The sensitivity section has a different role from the main paid evidence. The main evidence reports what happened in completed runs. Sensitivity analysis asks whether the same completed outcomes still support the cost-value story when reasonable evaluation choices change. We therefore treat sensitivity as a frontier diagnostic: each table or curve explains how BudgetFlow behaves when value, budget pressure, or serving cost changes.

7.1. Task Value Sensitivity

Task Value sensitivity tests whether the latest positive result depends on one narrow value profile. Table 6 re-scores the latest audited run under alternate Task Value profiles. BudgetFlow remains ahead of the best control under equal value, criticality value, compressed criticality, and expanded criticality. A value permutation diagnostic over 64 permutations gives BudgetFlow a positive margin in all samples, with median margin +1.00 TRV. The conclusion is that the latest




Figure 1. Multi-run main evidence. Each pair compares BudgetFlow with the best available control in that run on TRV per Dollar. The figure separates the repeated-run evidence from the sensitivity analysis in Section 7.

Value profile	BF TRV	Control TRV	Margin
equal	16.00	15.00	+1.00
criticality	18.00	17.00	+1.00
compressed	17.00	16.00	+1.00
expanded	20.00	19.00	+1.00

Table 6. Task Value sensitivity on the latest audited run.

KV profile	BF	Best control	BF margin
KV0	1.81	1.71	+0.10
KV50	3.26	3.14	+0.11
KV90	9.06	8.98	+0.09
KV98	14.08	14.28	-0.20
KV99	15.13	15.41	-0.29

Table 7. KV Cache Cost-Discount sensitivity. Values are TRV per Dollar.

frontier movement is visible across several value ledgers, including the plain resolved-count view where all resolved tasks have equal value.

7.2. Budget Sensitivity

Budget sensitivity is the closest analogue to a frontier plot. Figure 2 fixes observed outcomes and replays budget accounting at different caps. It therefore describes the observed cost-value frontier for this completed run. The strong-model-only policy is strongest at some low and mid caps, while BudgetFlow leads at the full cap in the paid run. The important reading is the frontier shape: the best policy changes with budget pressure, and BudgetFlow reaches the top frontier point when the full shared budget is available.

7.3. KV Cache Cost-Discount Sensitivity

KV Cache Cost-Discount sensitivity tests a serving-cost question. Modern serving substrates can make repeated input context cheaper; this changes effective spend while leaving verified outcomes unchanged. Table 7 re-costs repeated input tokens under KV-cache discounts. BudgetFlow leads or stays close through KV90. At KV98 and KV99, the learned task-router has the highest TRV per Dollar, while BudgetFlow keeps a higher TRV than the learned task-router in the fixed-outcome table.

This pattern is informative rather than contradictory. Budget-

Flow gains part of its efficiency by allocating scarce budget across tasks. When cache discounts become extreme, the cost penalty of repeated or stronger model use shrinks, and value-blind controls can become more competitive on TRV per Dollar. The result makes the serving boundary explicit: BudgetFlow is most useful under meaningful budget pressure, and serving optimizations can shift which policy lies on the efficiency frontier.

8. Discussion and Future Work

Formulation-level generalization. The SWE-bench experiments instantiate the formulation in a verified coding domain. The formulation applies to task domains where tasks expose Task Value, Estimated Token Demand, Model Fit, and a Verifier. Coding tasks use tests as the Verifier. Support tasks can use policy checks and service outcomes. Spreadsheet tasks can use cell-level checks or rule engines. Research and document-production tasks can use human review or domain checks. The calibration of each interface is domain-specific; the allocation objective and budget ledger are shared.

Boundary-aware reporting. The four paid runs show different frontier leaders. This is the intended evidence shape for Batch-Level Task Budget Allocation. A value-




Figure 2. Budget sensitivity diagnostic on the latest audited run. Outcomes are fixed; budget accounting is replayed at alternate caps. The dashed line is the observed frontier; BudgetFlow is emphasized in red, while controls are shown as lighter diagnostic curves.

aware policy can improve TRV and TRV per Dollar, a strong model can define the best boundary in favorable conditions, and a learned task-router can be a strong control when value placement is less important. Reporting the frontier makes these cases visible.

Future work. Fine-grained task execution control. BudgetFlow can move from task-level allocation to finer-grained execution control inside a task. The current paper assigns model opportunities at the task level. **A richer policy could decide which task stage or segment deserves the strong model, how much repair runway to preserve, and when stop/defer decisions should be made inside a long attempt.** This direction should be evaluated against task-level controls because finer granularity can improve allocation only when the stage signal is reliable enough to justify the added switching cost.

Continual policy learning with auditability. Completed batches produce cost traces, routing outcomes, verifier results, and frontier positions. **A future policy can use those records to update future caps, route choices, and stop/continue decisions.** The important systems constraint is that each learned update should remain inspectable: the next allocation should be explainable through the same value ledger, demand estimate, model-fit evidence, and verifier-settled outcome history.

Serving-aware and multi-tenant BudgetFlow. Serving systems such as vLLM, SGLang, and Dynamo optimize execution substrates: batching, cache locality, scheduling, and inference throughput. BudgetFlow operates at the task-budget layer above that substrate. A serving-aware Budget-

Flow could pass task value, remaining budget, estimated demand, and outcome history into scheduling hints such as priority, continuation, model opportunity, and cache-sticky execution. The same interface can then extend from one budget owner to multiple teams, users, or services that share an agent execution substrate with distinct budgets and priorities. **This turns Batch-Level Task Budget Allocation into a candidate control plane for multi-tenant agent workloads.**

9. Conclusion

This paper formulates Batch-Level Task Budget Allocation for multi-step agent tasks under a shared hard budget. The core question is which task should receive the next scarce model opportunity when tasks have different value, different expected demand, and different model-tier fit. BudgetFlow answers this question with a simple auditable interface: Task Value, Estimated Token Demand, Model Fit, a shared budget ledger, and a Verifier-settled outcome.

The SWE-bench experiments instantiate this interface in a verified coding domain. In the latest audited 5-policy run, BudgetFlow improves Resolved Rate, TRV, and TRV per Dollar. Across the broader paid evidence, different policies occupy the frontier under different operating conditions, which is precisely why the paper reports cost-value curves rather than a single point estimate. The sensitivity analyses show how the frontier responds to value profiles, budget caps, and serving cost assumptions.

The broader value of BudgetFlow is the formulation. Coding tasks provide a strict verifier and measurable spend, but

the same batch-budget structure appears whenever many agent tasks compete for one resource pool and produce outcomes with unequal value. **By separating allocation signals from outcome settlement, BudgetFlow offers a portable way to reason about verified value under fixed resource constraints, from SWE-bench-style coding to future serving-aware and multi-tenant agent systems.**

References

- Dekoninck, J., Baader, M., and Vechev, M. A unified approach to routing and cascading for LLMs. In *International Conference on Machine Learning*, 2025. URL <https://arxiv.org/abs/2410.10347>.
- Guo, D., Wu, J., and Yiu, S. M. RouteNLP: Closed-loop LLM routing with conformal cascading and distillation co-optimization. In *ACL Industry Track*, 2026. URL <https://arxiv.org/abs/2604.23577>.
- Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., and Narasimhan, K. SWE-bench: Can language models resolve real-world github issues? In *International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2310.06770>.
- Kotte, V. UCCI: Calibrated uncertainty for cost-optimal LLM cascade routing. *arXiv preprint arXiv:2605.18796*, 2026. URL <https://arxiv.org/abs/2605.18796>.
- Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., Gonzalez, J. E., Zhang, H., and Stoica, I. Efficient memory management for large language model serving with pagedattention. In *ACM Symposium on Operating Systems Principles*, 2023. URL <https://arxiv.org/abs/2309.06180>.
- Liu, H., Tian, C., An, N., Wang, Z., Lu, P., Yu, C., and Qi, Q. Budget-constrained agentic large language models: Intention-based planning for costly tool use. *arXiv preprint arXiv:2602.11541*, 2026. URL <https://arxiv.org/abs/2602.11541>.
- NVIDIA. NVIDIA Dynamo: Agentic inference. <https://docs.nvidia.com/dynamo/digest/agentic-inference>, 2026. Accessed 2026-07-06.
- Ong, I., Almahairi, A., Wu, V., Chiang, W.-L., Wu, T., Gonzalez, J. E., Kadous, M. W., and Stoica, I. RouteLLM: Learning to route LLMs with preference data. In *International Conference on Learning Representations*, 2025. URL <https://arxiv.org/abs/2406.18665>.
- OpenSquilla contributors. OpenSquilla: Token-efficient ai agent. <https://github.com/opensquilla/opensquilla>, 2026. Accessed 2026-07-06.
- Zheng, L., Yin, L., Xie, Z., Sun, C., Huang, J., Yu, C. H., Cao, S., Kozyrakis, C., Stoica, I., Gonzalez, J. E., Barrett, C., and Sheng, Y. SGLang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems*, 2024. URL <https://arxiv.org/abs/2312.07104>.
- Zheng, M., Han, K., Li, B., Xu, H., Tian, Y., He, W., Zhou, H., Guo, J., Hu, H., Ma, L., Xu, C., Dai, G., Xia, L., Wei, Y., Wang, Y., and Wang, Y. Claw-SWE-Bench: A benchmark for evaluating openclaw-style agent harnesses on coding tasks. *arXiv preprint arXiv:2606.12344*, 2026. URL <https://arxiv.org/abs/2606.12344>.